

Parallel Computing



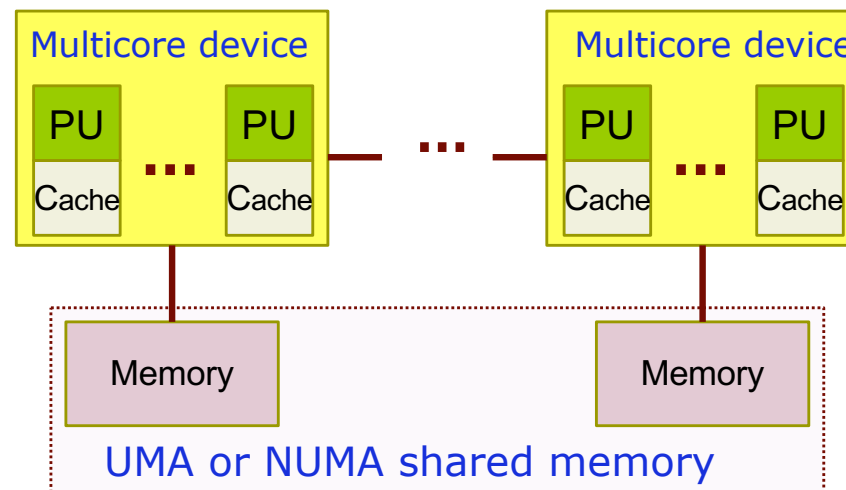
João Luís Ferreira Sobral
www.di.uminho.pt/~jls
jls@...

Web: Elearning

Explicit parallel computing (1)

Shared Memory (SM) parallel systems

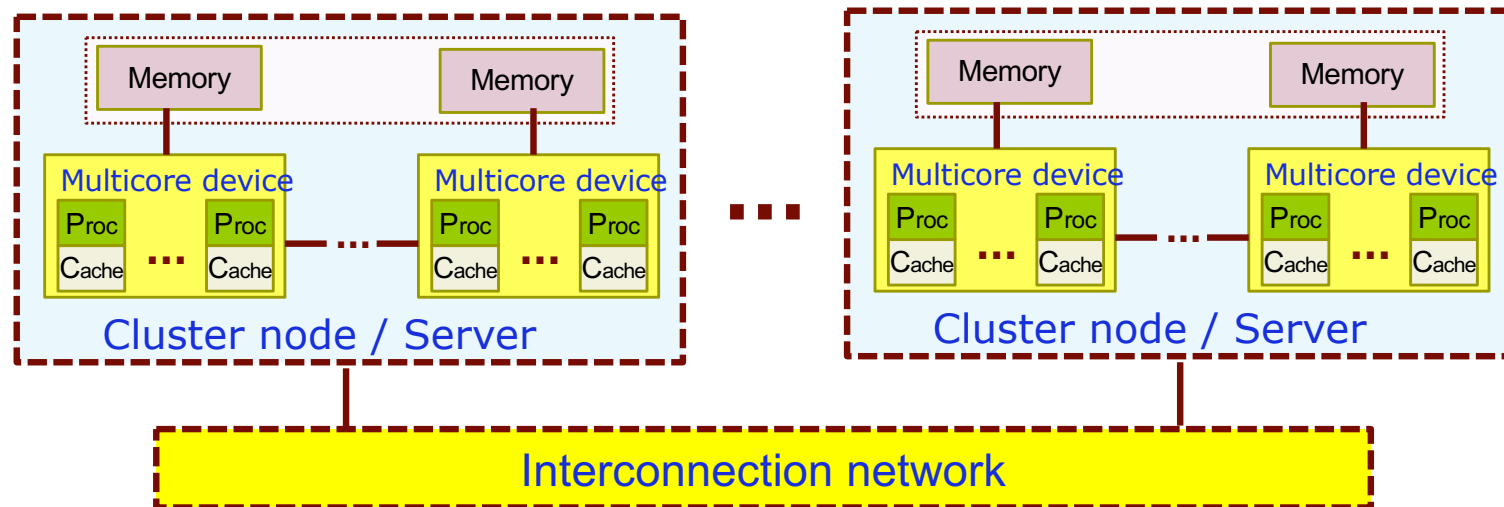
- parallelism on single or multiple devices (*same motherboard*)
 - single physical memory address space
 - each core (PU) can support multiple threads (SMT)
 - **memory bandwidth is shared by all cores**
 - **coherence of data in multiple caches?**



Explicit parallel computing (2)

Distributed memory parallel systems

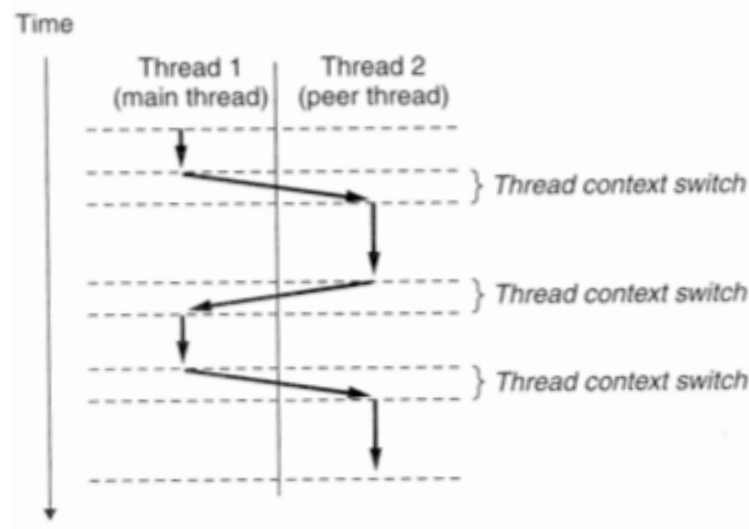
- on multiple boards (*or multiple nodes/servers*)
 - each node with its private memory space
 - **memory bandwidth is proportional to the number of nodes**




Thread Scheduling

(on shared memory parallel systems)

Review: OS thread scheduling and context switch



Superscalar
PU (4-way)



Thread scheduling at the Operating System (OS) level:

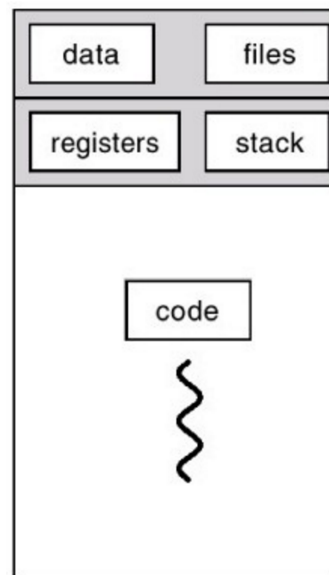
1. Interrupts execution of **Thread1** & saves **Thread1** context information (e.g., PU registers)
2. Selects the next thread to execute (**Thread2**)
3. OS restores the state of **Thread2** and gives control to it
4. **Thread2** continues execution from its previous execution point (saved Program Counter register)

Programming model: Process vs Threads

□ Processes

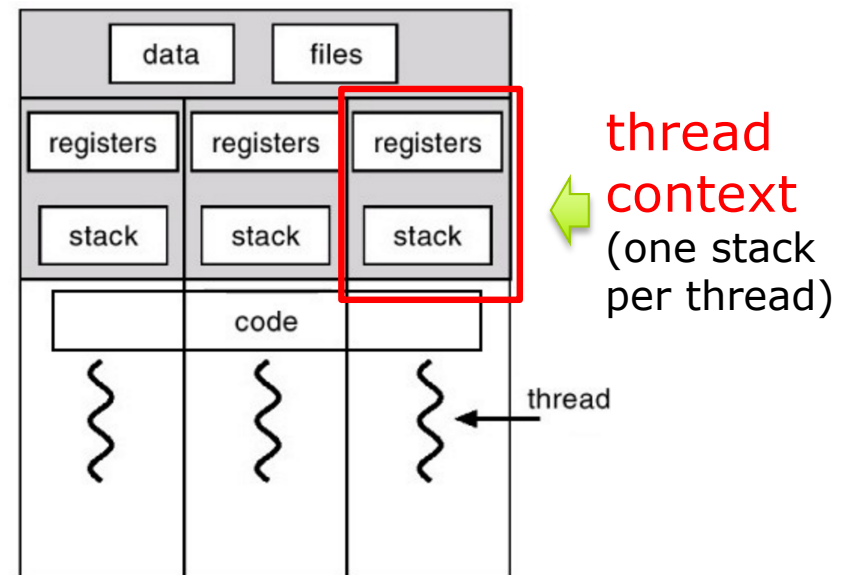
- Used for unrelated tasks
 - (e.g., a program)
- Own address space
 - Address space is protected from other process
- Switching at the kernel (OS) level

Every process has at least one thread

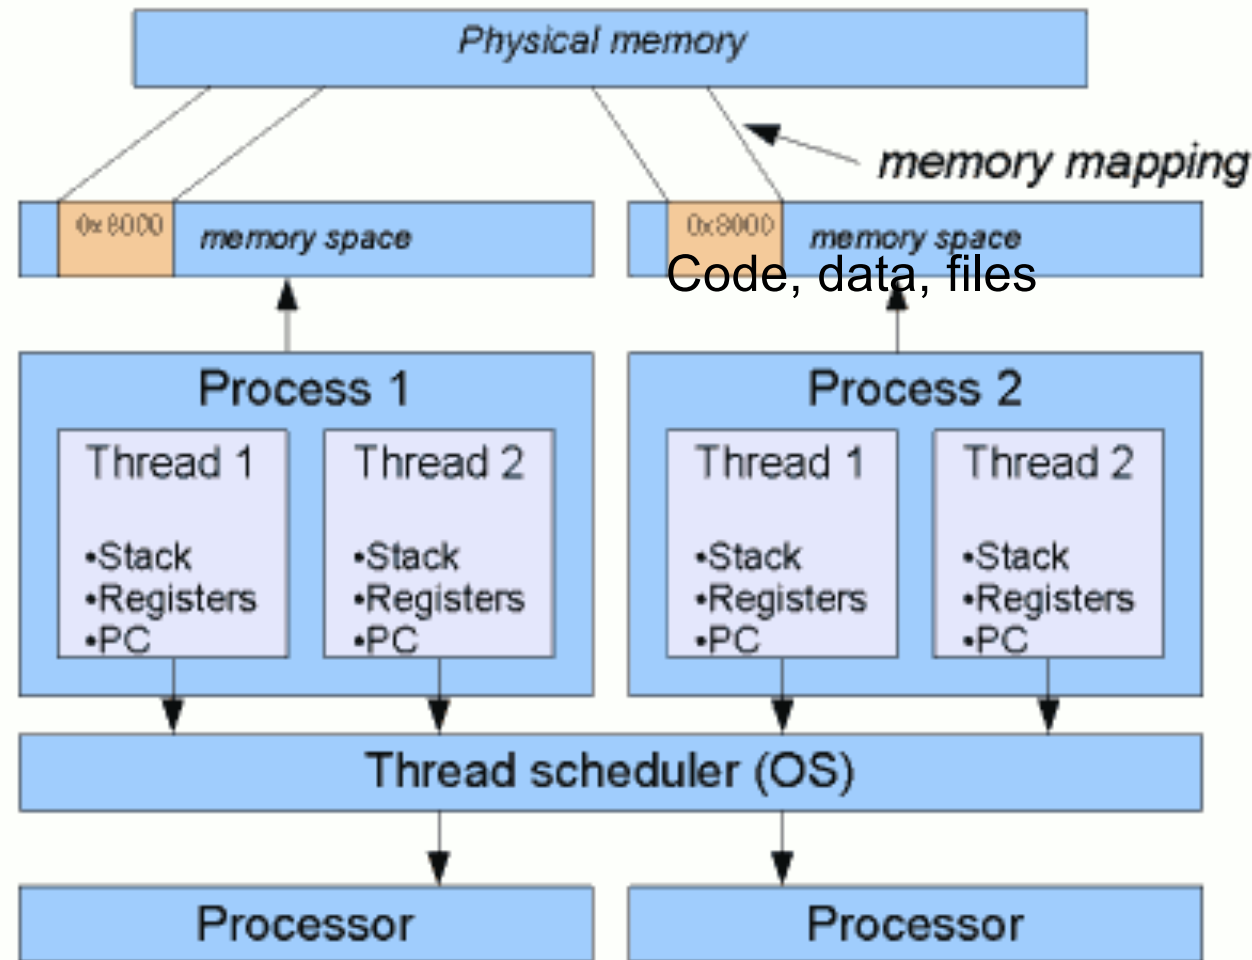


□ Threads

- Are part from the same job
- Share address space, code, data and files
- Switching at the user or kernel level



Thread vs Process




← process memory space is shared by all threads of a process

Threads from all processes are scheduled into available PUs

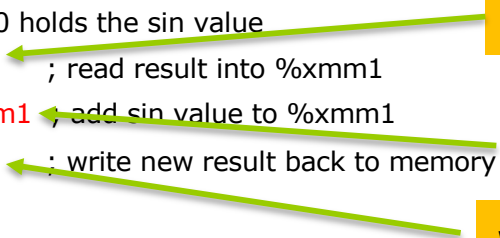
Data races

- A data race **can happen** when two or more threads access (write!) to a **shared memory position**

```
5 int main(){
6     double result={0};
7
8     #pragma omp parallel for shared(result)
9     for(int i=0; i<1000000;i++) {
10         result+=sin(i);
11     }
12     printf("%f",result);
13 }
```



```
.L4:
...
call sin      ; after sin call, xmm0 holds the sin value
vmovsd 8(%rsp,r12), %xmm1 ; read result into %xmm1
vaddsd %xmm0, %xmm1, %xmm1 ; add sin value to %xmm1
vmovsd %xmm1, 8(%rsp,r12) ; write new result back to memory
...
cmpl %ebx, %ebp ; i<1000000?
jne .L4
```



Thread 0

read result

add sin

write result

Thread 1

read result

add sin

write result

Process/Thread vs Tasks

- **Task:** sequence of instructions
- **Thread/process:** execution context for a task
- **Processor/core:** hardware that runs a thread/process

In Java

- Runnable object
- Thread
- Processor core

